

ORGANISASI DAN ARSITEKTUR KOMPUTER

Pipeline processing

1

PIPELINE DASAR

- What is pipelining
- Clock & Latches
- Contoh 5 Stage Pipeline
- Load/Store & RISC/CISC
- Hazard
- Examples of Hazards

2

GENERAL CONCEPTS

- Pipelining merupakan teknik yang membagi task kedalam sejumlah subtask yang perlu dilakukan dalam sebuah *sequence*.
- Setiap subtask dikerjakan oleh sebuah fungsional unit. Unit-unit terhubung secara serial dan semuanya beroperasi secara simultan.
- Penggunaan pipelining untuk memperbaiki performa.

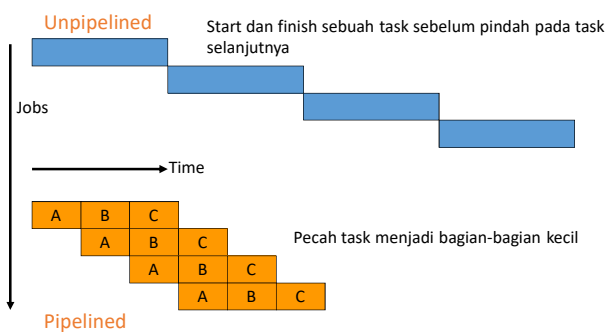
3

GENERAL CONCEPTS

- Teknik pipeline ini dapat diterapkan pada berbagai tingkatan dalam sistem komputer.
- Bisa pada level yang tinggi, misalnya program aplikasi, sampai pada tingkat yang rendah, seperti pada instruksi yang dijalankan oleh microprocessor.

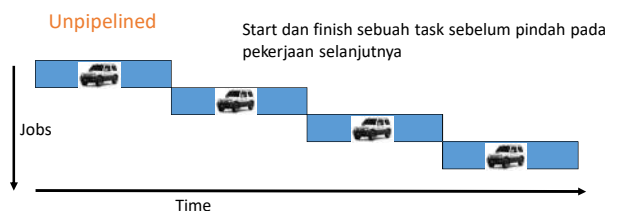
4

The Assembly Line



5

ILUSTRASI

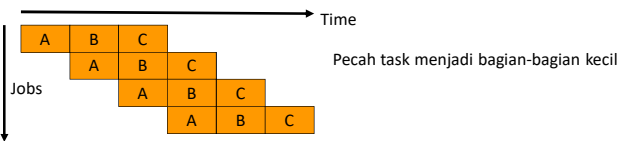


- Misal untuk 1 mobil butuh 24 jam
- Tidak ada paralelisme
- Troughput → 1mobil/24jam

6

ILUSTRASI

Pipelined



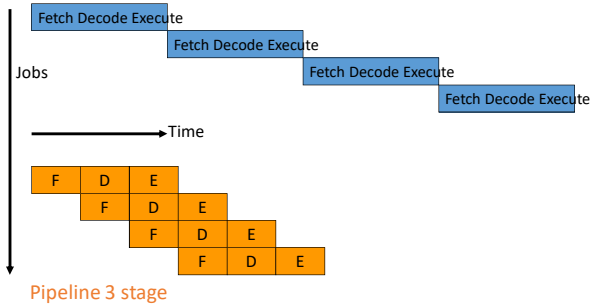
Misal:
A → Engine
B → Body
C → Paint

- Dengan paralelisme meningkatkan hasil
- Troughput → 1mobil/8jam

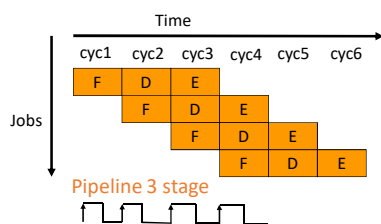
Waktu per stage = 8 jam

7

Unpipelined

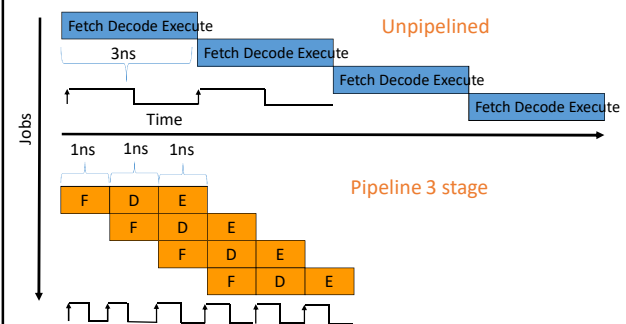


8



- Setiap clock naik → sebuah instruksi selesai

9



Unpipelined clock speed= $1/3\text{ns} = 333\text{MHz}$

Pipelined clock speed= $1/1\text{ns} = 1\text{GHz}$

10

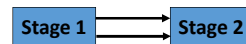
EFEK

Maka dengan pipelining:

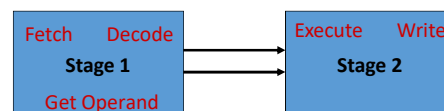
- Waktu per instruksi naik
- Jumlah siklus per instruksi naik (perhatikan peningkatan kecepatan clock)
- Total waktu eksekusi turun
- Jumlah pipeline stage = peningkatan kecepatan clock

11

CLOCK & LATCHES

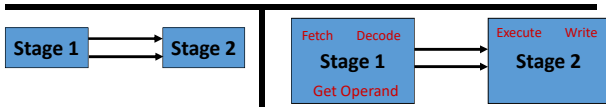


Misal:



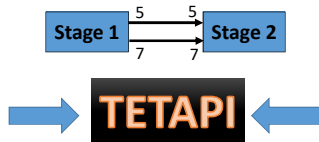
12

CLOCK & LATCHES



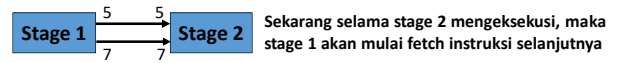
Misal: I₁: ADD R1+R2 → R3 ;asumsi R1=5 R2=7

- Jadi selama stage 1 kita mengambil nilai di R1 dan R2
- Di akhir stage 1 → output stage 1 jadi input untuk stage 2

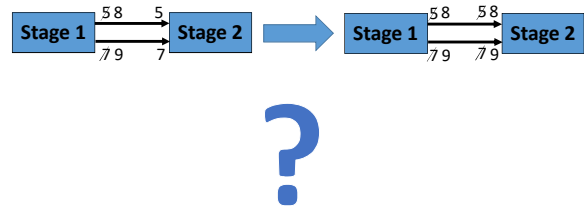


13

CLOCK & LATCHES



- I₁ : ADD R1+R2 → R3 ;asumsi R1=5 R2=7
- I₂ : ADD R4+R5 → R6 ;asumsi R4=8 R5=9

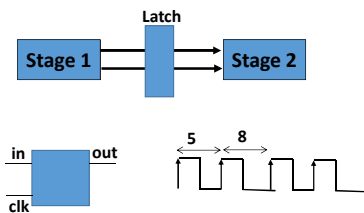


14

CLOCK & LATCHES

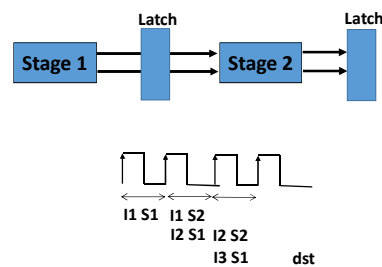
Butuh keadaan dimana saat stage 2 mengeksekusi data input-nya tidak berubah
SOLUSI:

- Memisahkan stage 1 dan stage 2 dengan menggunakan Latch.



15

CLOCK & LATCHES

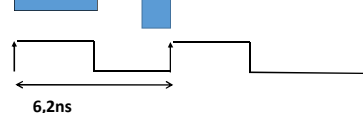


16

CLOCK & LATCHES

Overhead

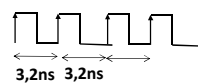
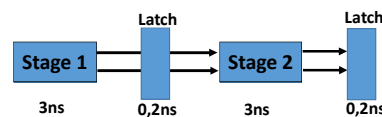
Misal: tanpa pipeline 1 instruksi menghabiskan 6ns dan untuk menyimpan hasil ke latch 0,2ns



Cycle time = 6,2ns
Clock speed = 1/6,2ns = 160MHz

17

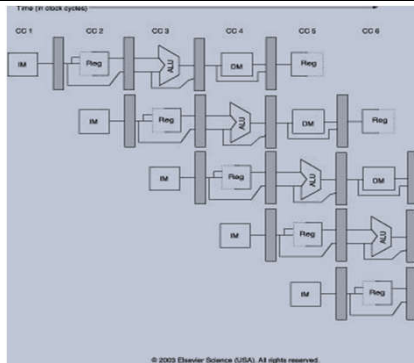
CLOCK & LATCHES



Cycle time = 3,2ns
Clock speed = 1/3,2ns = 312,5MHz

18

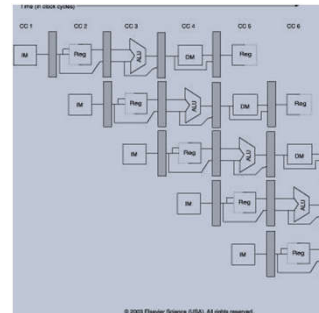
A 5-Stage Pipeline



DETAIL SETIAP STAGE

A 5-Stage Pipeline

Use the PC to access the I-cache and increment PC by 4



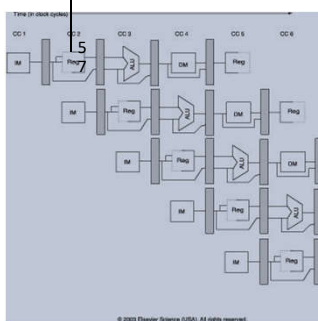
- PC
- Nilai PC bertambah
- Latch

PC

20

A 5-Stage Pipeline

Read registers, compare registers, compute branch target; for now, assume branches take 2 cyc (there is enough work that branches can easily take more)



Dec. RR

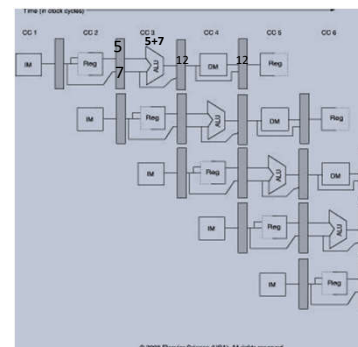
Misal:
ADD R1+R2

Misal jika ada instruksi lompat:
BRZ R3,[24]

21

A 5-Stage Pipeline

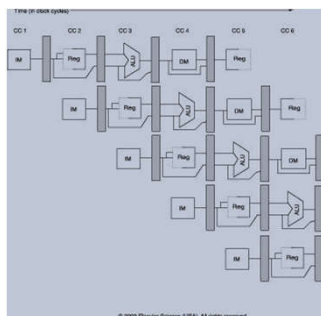
ALU computation, effective address computation for load/store



22

A 5-Stage Pipeline

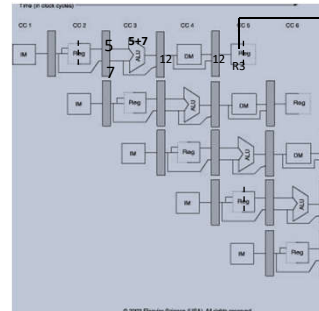
Memory access to/from data cache, stores finish in 4 cycles



23

A 5-Stage Pipeline

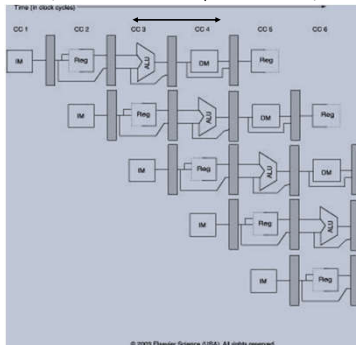
Write result of ALU computation or load into register file



24

A 5-Stage Pipeline

ALU computation, effective address computation for load/store



25

RISC/ CISC Load/Stores

- RISC
 - Reduce instruction set computer
 - Setiap instruksi sederhana

Misal:

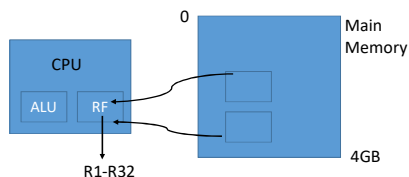
ADD R1,R2 → R3

OR R1,R2 → R3

LD []→R4

ST []←R5

26



Misal:

Kita ingin menyalin nilai dari Main memory ke RF (register File) R1 dan R2 dan menjumlahkannya dan simpan ke R3 lalu disimpan ke alamat memori yang ditunjuk.

Jadi:

LD []→R1

LD []→R2

Add R1,R2 →R3

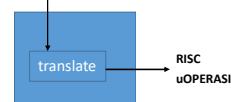
ST []←R3

27

- CISC (Complex Instruction Set Computer)
- 1 sequence instruksi sama dengan beberapa instruksi sederhana

- MUL R1,R2→R3
 - ADD R3,R4 →R4
- • MAC R1,R2,R4

Intel x86 → CISC



28

Misal:

- ALU ADD/OR/SUB R1,R2→R3

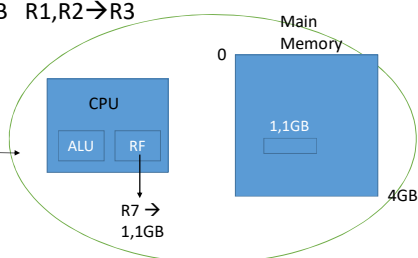
- LD []→R4

- ST []←R5

- LD [R7]→R4

- LD 8[R7]→R4

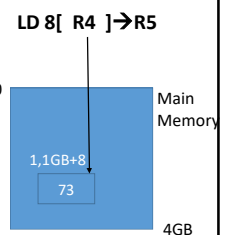
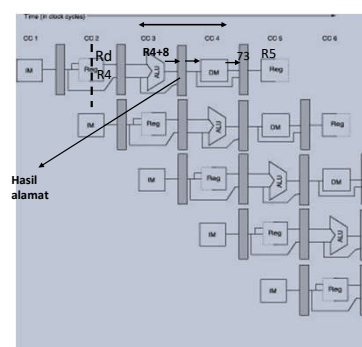
- ST 24[R7]←R5



29

A 5-Stage Pipeline

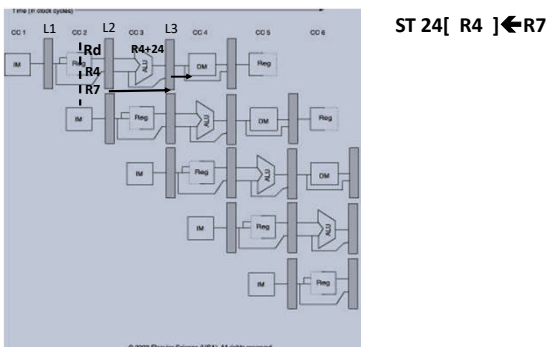
ALU computation, effective address computation for load/store



30

A 5-Stage Pipeline

ALU computation, effective address computation for load/store



31

TERIMA KASIH

32

Hazards

- Hazard adalah keadaan yang dapat menimbulkan tunda (delay, stall) pada pipeline.
- Pada keadaan stall, pipeline tidak menghasilkan output sehingga peningkatan keluaran ideal tidak dapat dicapai.

33

Hazards

- Operasi pipeline akan stall jika salah satu unit atau stage membutuhkan lebih banyak waktu untuk melakukan fungsinya dan memaksa stage lainnya untuk idle.
- Situasi seperti itu disebut pipeline bubble (pipeline hazards)

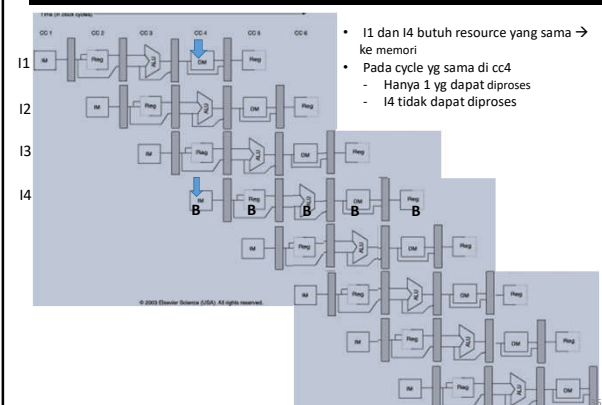
34

Hazards

- Structural hazards: different instructions in different stages (or the same stage) conflicting for the same resource
- Data hazards: an instruction cannot continue because it needs a value that has not yet been generated by an earlier instruction
- Control hazard: fetch cannot continue because it does not know the outcome of an earlier branch – special case of a data hazard – separate category because they are treated in different ways

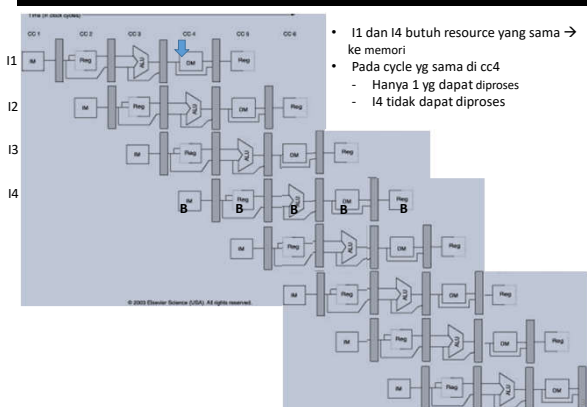
35

A 5-Stage Pipeline

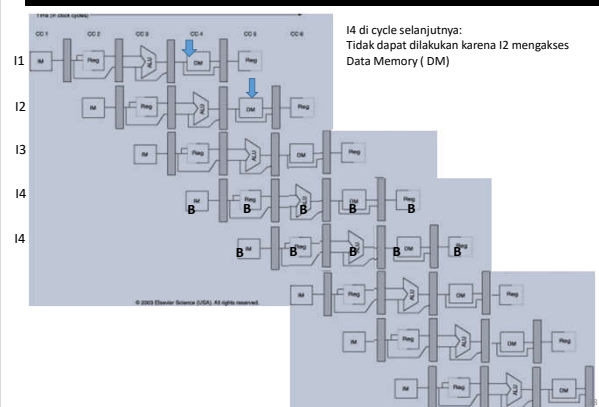


36

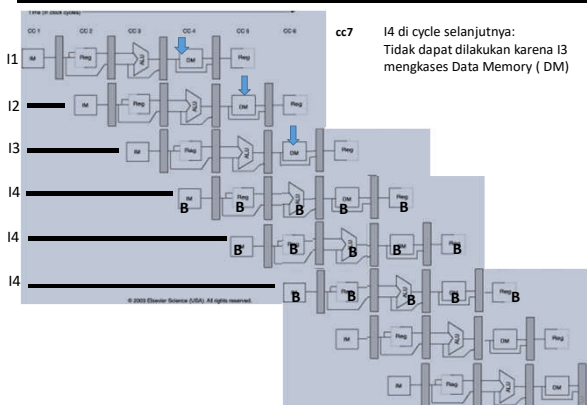
A 5-Stage Pipeline



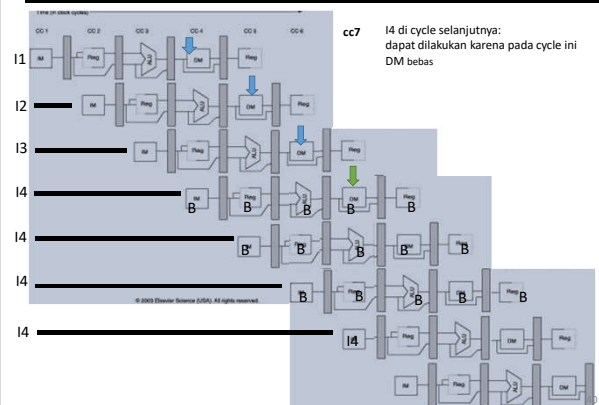
A 5-Stage Pipeline



A 5-Stage Pipeline



A 5-Stage Pipeline

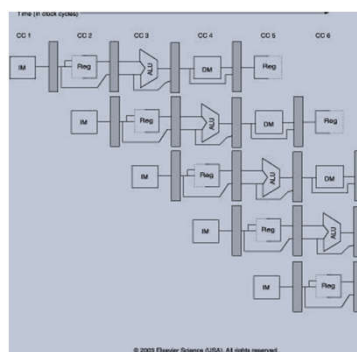


Solusi:

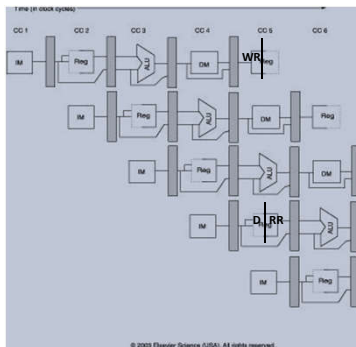
Membagi memori menjadi dua yaitu:

- memori instruksi dan
- memori data

Masalah di cc5 bagian register file



Masalah di cc5 bagian register file



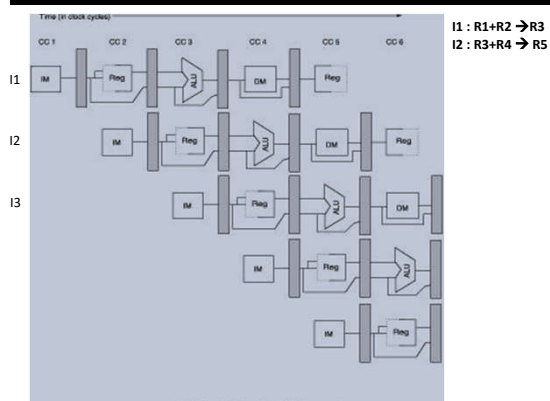
43

Structural Hazards

- Example: a unified instruction and data cache → stage 4 (MEM) and stage 1 (IF) can never coincide
- The later instruction and all its successors are delayed until a cycle is found when the resource is free → these are pipeline bubbles
- Structural hazards are easy to eliminate – increase the number of resources (for example, implement a separate instruction and data cache)

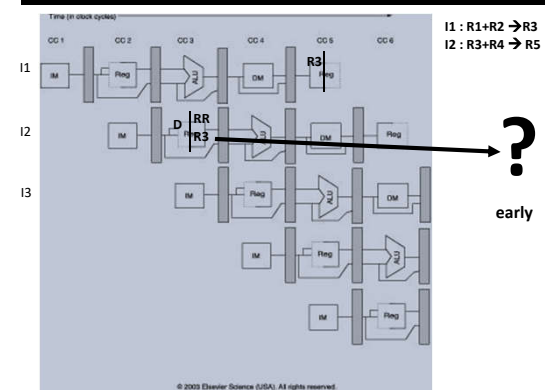
44

Data Hazards



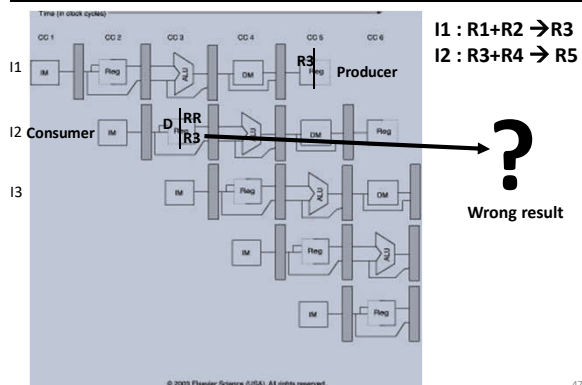
45

Data Hazards



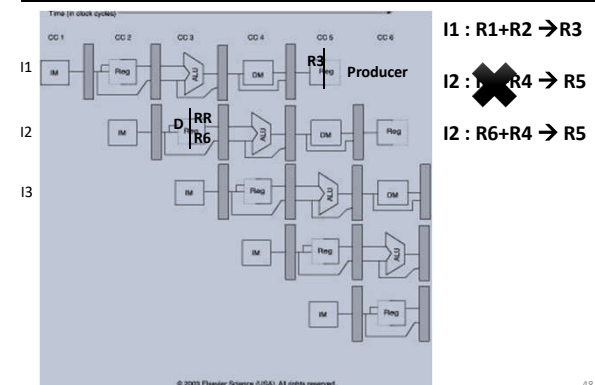
46

Data Hazards



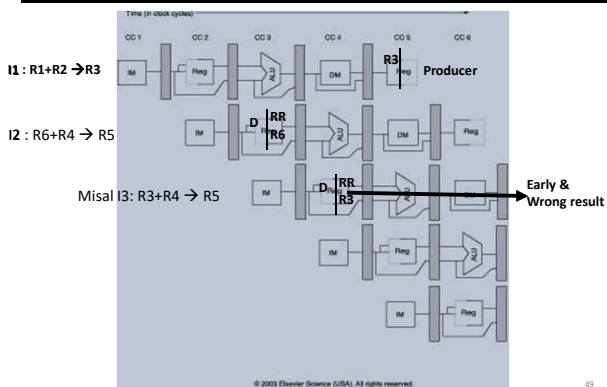
47

Data Hazards

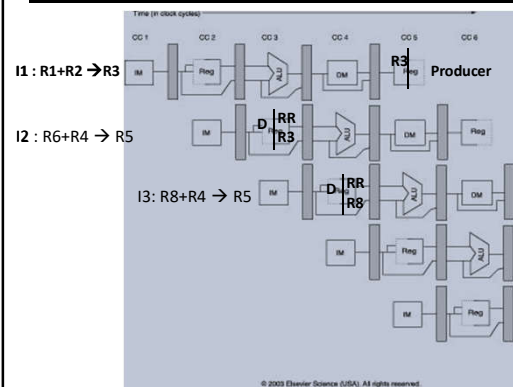


48

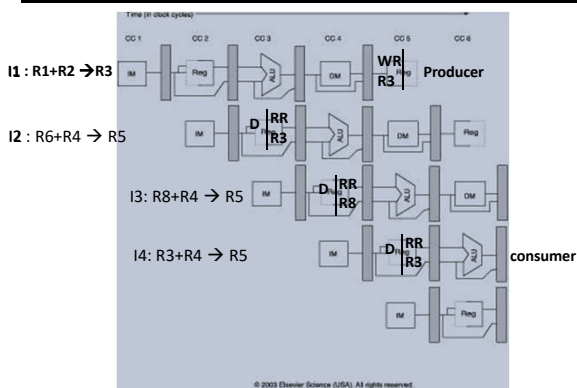
Data Hazards



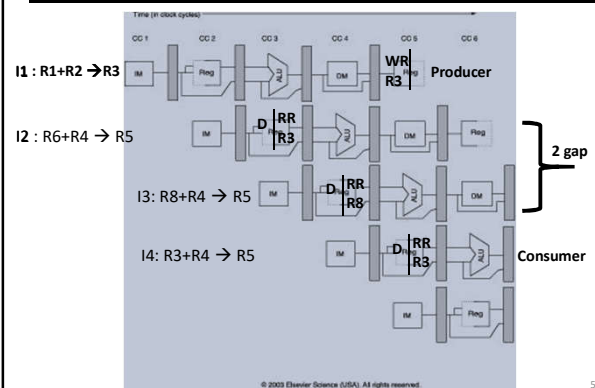
Data Hazards



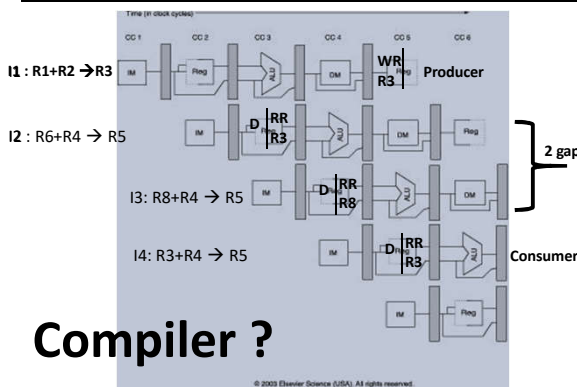
Data Hazards



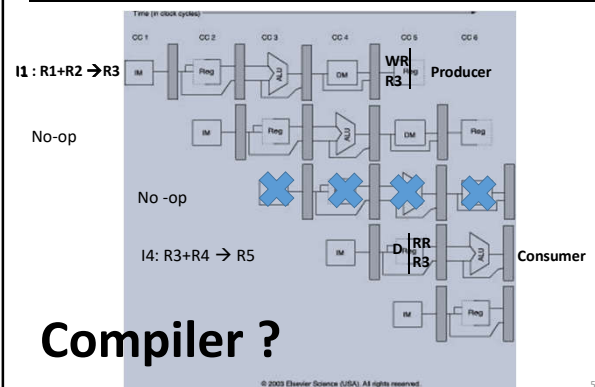
Data Hazards



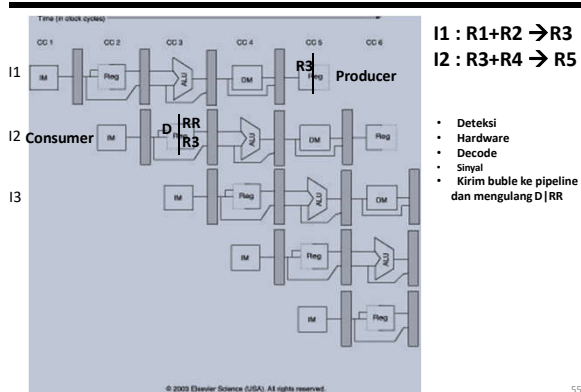
Data Hazards



Data Hazards

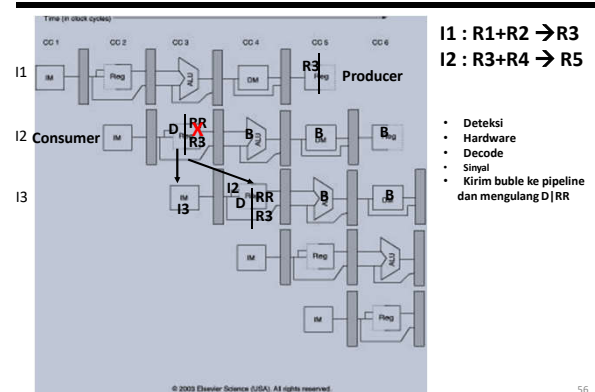


Data Hazards



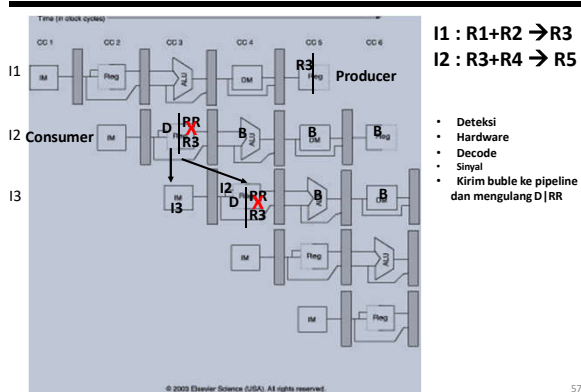
55

Data Hazards



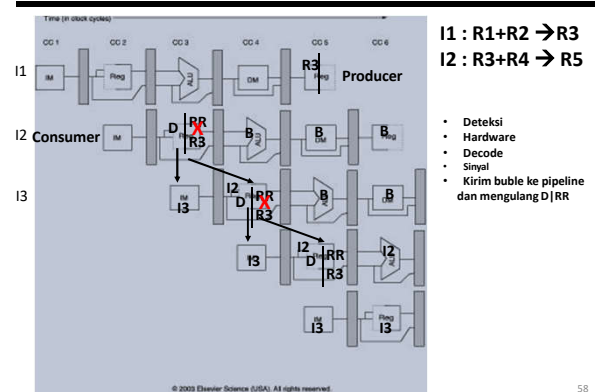
56

Data Hazards



57

Data Hazards

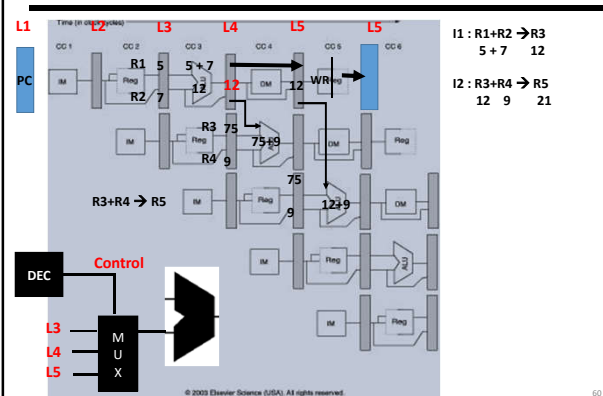


58



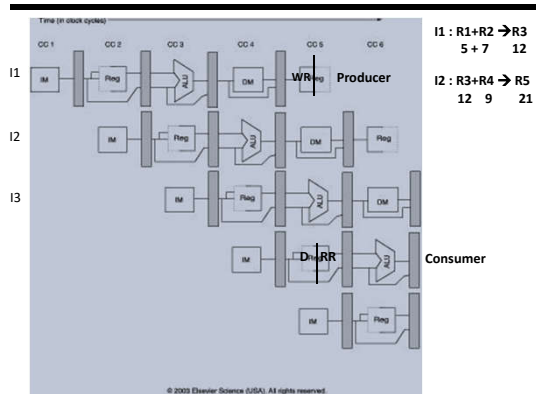
59

A 5-Stage Pipeline



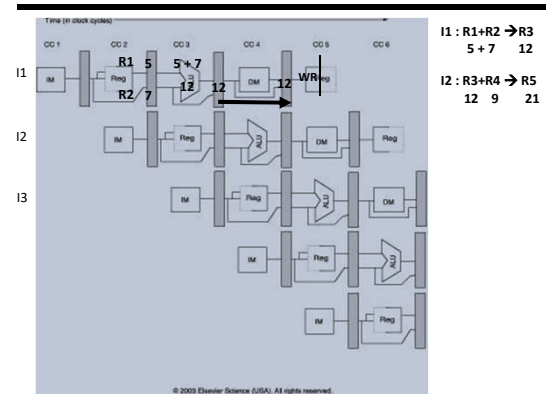
60

A 5-Stage Pipeline



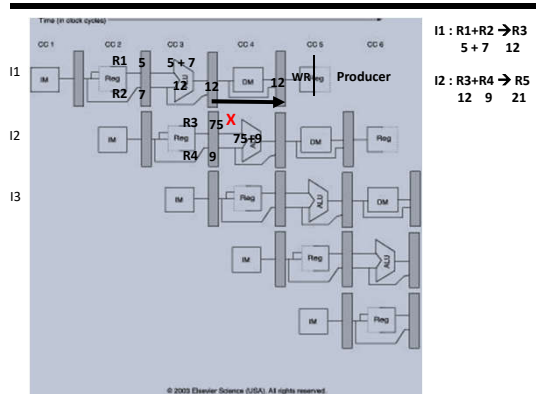
G1

A 5-Stage Pipeline



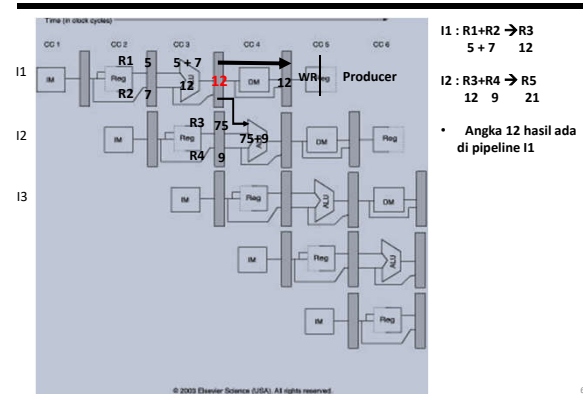
G2

A 5-Stage Pipeline



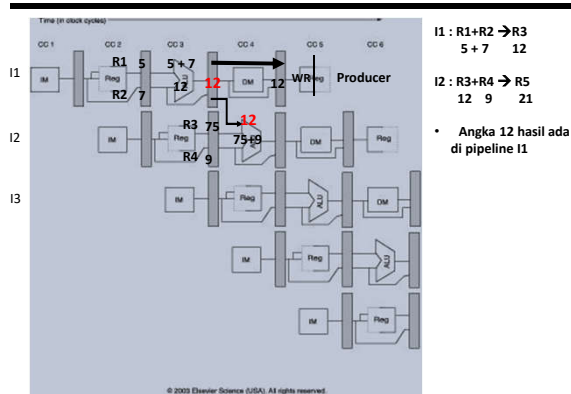
G3

A 5-Stage Pipeline



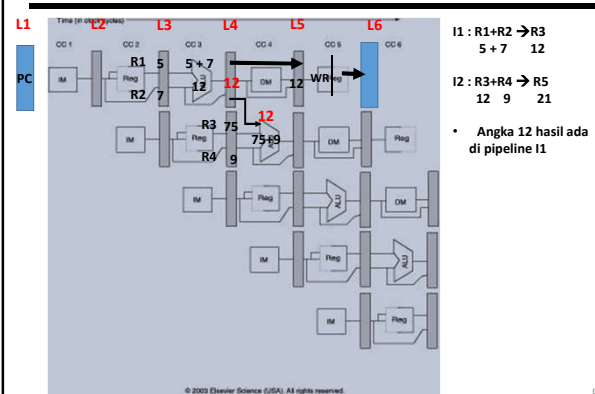
G4

A 5-Stage Pipeline



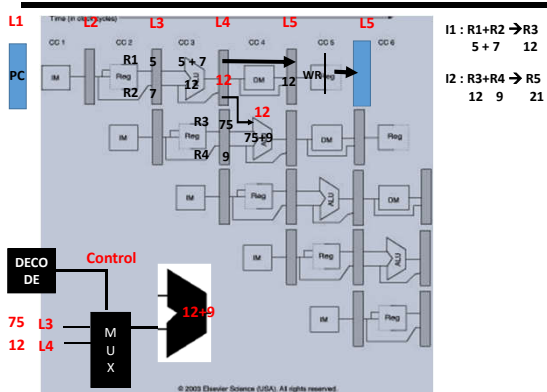
G5

A 5-Stage Pipeline



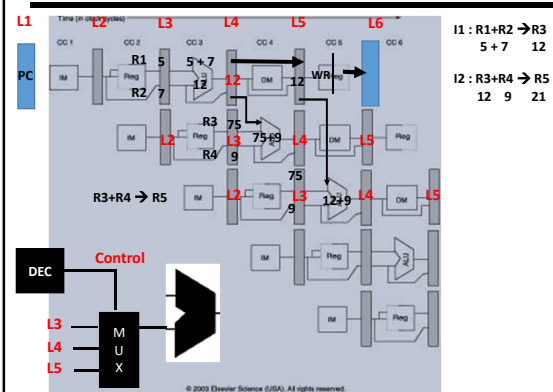
G6

A 5-Stage Pipeline



67

A 5-Stage Pipeline



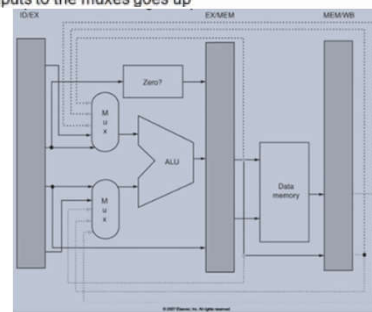
68



69

Pipeline Implementation

- Signals for the muxes have to be generated – some of this can happen during ID
- Need look-up tables to identify situations that merit bypassing/stalling – the number of inputs to the muxes goes up

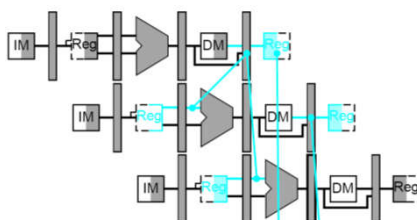


70

Example

add R1, R2, R3
 $(R2+R3 \rightarrow R1)$

lw R4, 8(R1)



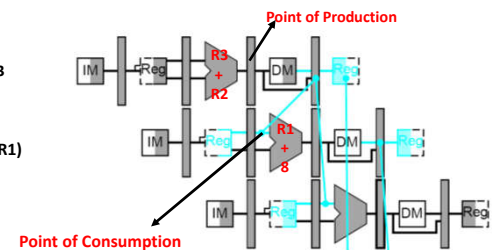
- Point of Production
- Point of Consumption

71

Example

add R1, R2, R3
 $(R2+R3 \rightarrow R1)$

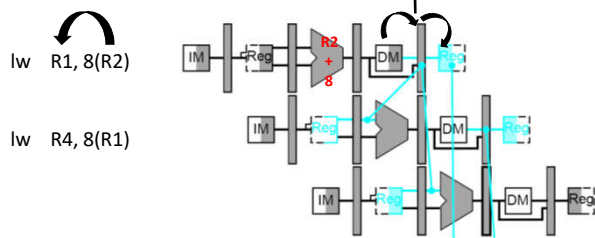
lw R4, 8(R1)



Point of Consumption

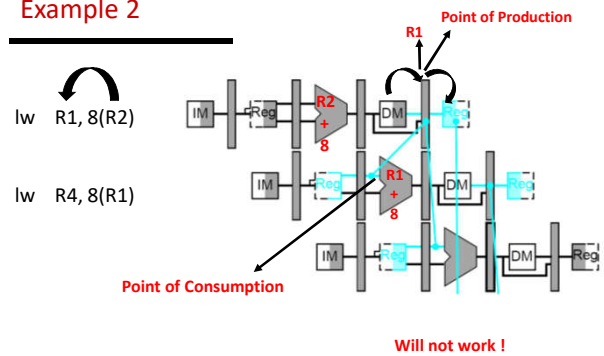
72

Example 2



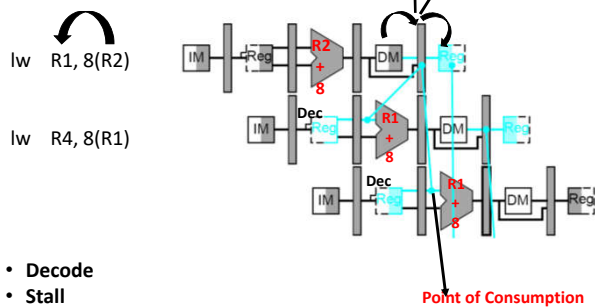
73

Example 2



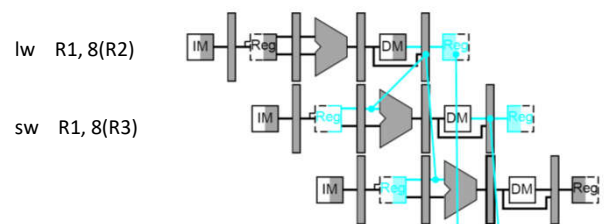
74

Example 2



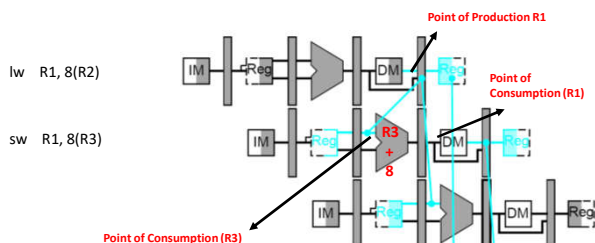
75

Example 3



76

Example 3



77

Summary

- For the 5-stage pipeline, bypassing can eliminate delays between the following example pairs of instructions:
 add/sub R1, R2, R3
 add/sub/lw/sw R4, R1, R5

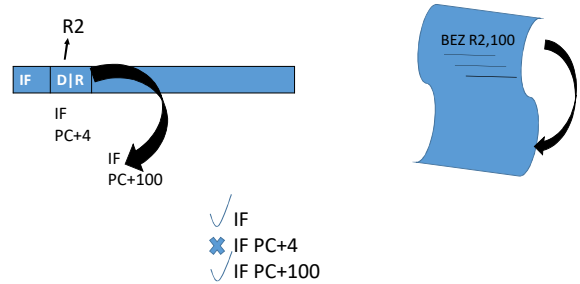
 lw R1, 8(R2)
 sw R1, 4(R3)
- The following pairs of instructions will have intermediate stalls:
 lw R1, 8(R2)
 add/sub/lw R3, R1, R4 or sw R3, 8(R1)

78



79

Control Hazards



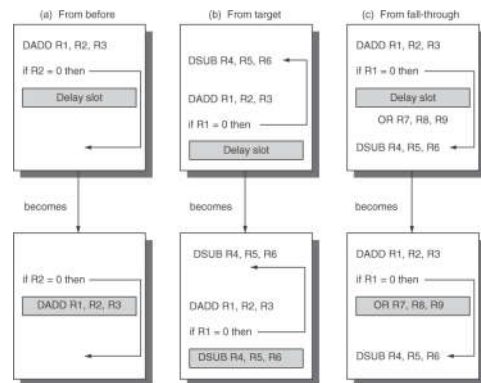
80

Control Hazards

- Simple techniques to handle control hazard stalls:
 - for every branch, introduce a stall cycle (note: every 6th instruction is a branch!)
 - assume the branch is not taken and start fetching the next instruction – if the branch is taken, need **hardware** to cancel the effect of the wrong-path instruction
 - Predict the next PC and fetch that instruction
 - fetch the next instruction (branch delay slot) and execute it anyway – if the instruction turns out to be on the correct path, useful work was done – if the instruction turns out to be on the wrong path, hopefully program state is not lost

81

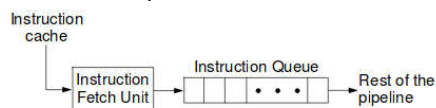
Branch Delay Slots



82

Instruction Fetch Units and Instruction Queues

- Most processors employ sophisticated fetch units that fetch instructions before they are needed and store them in a queue.



- The fetch unit also has the ability to recognize branch instructions and to generate the target address.

83

Instruction Fetch Units and Instruction Queues

- Penalty produced by unconditional branches can be drastically reduced: the fetch unit computes the target address and continues to fetch instructions from that address, which are sent to the queue.
- The rest of the pipeline gets a continuous stream of instructions, without stalling.

84

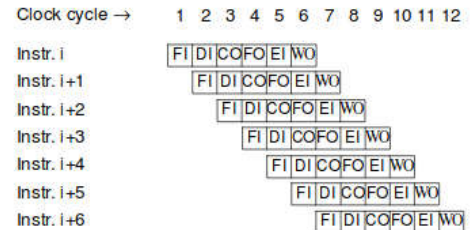
Instruction Fetch Units and Instruction Queues

- The rate at which instructions can be read (from the instruction cache) must be sufficiently high to avoid an empty queue.
- With conditional branches penalties can not be avoided.
- The branch condition, which usually depends on the result of the preceding instruction, has to be known in order to determine the following instruction.

85

CONTROL HAZARD (6 stage pipeline example)

- FI: fetch instruction
- DI: decode instruction
- CO: calculate operand address
- FO: fetch operand
- EI: execute instruction
- WO: write operand



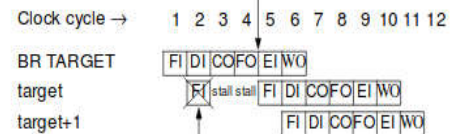
86

Unconditional branch



87

After the FO stage of the branch instruction the address of the target is known and it can be fetched



The instruction following the branch is fetched; before the DI is finished it is not known that a branch is executed. Later the fetched instruction is discarded

Penalty: 3 cycles

88

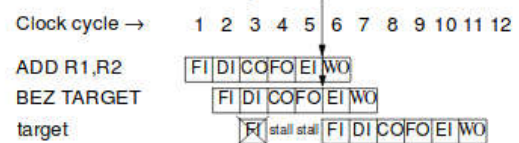
Conditional branch



89

Branch is taken

At this moment, both the condition (set by ADD) and the target address are known.



Penalty: 3 cycles

90

Branch not taken

At this moment the condition is known and instr+1 can go on.



Penalty: 2 cycles

With conditional branch we have a penalty even if the branch has *not* been taken. This is because we have to wait until the branch condition is available.

91

DELAYED BRANCHING

- The idea with delayed branching is to let the CPU do some useful work during some of the cycles which are shown above to be stalled.
- With delayed branching the CPU **always** executes the instruction that immediately follows after the branch and only then alters (if necessary) the sequence of execution. The instruction after the branch is said to be in the *branch delay slot*.

92

This is what the programmer has written

This instruction does not influence any of the instructions which follow until the branch; it also doesn't influence the outcome of the branch.

This instruction should be executed only if the branch is not taken.

```

→ MUL R3,R4  R3 ← R3*R4
  SUB #1,R2  R2 ← R2-1
  ADD R1,R2  R1 ← R1+R2
  BEZ TAR    branch if zero
→ MOVE #10,R1 R1 ← 10
  -----
  TAR        -----
  
```

- The **compiler** (assembler) has to find an instruction which can be moved from its original place into the branch delay slot after the branch and which will be executed regardless of the outcome of the branch.

93

This is what the compiler (assembler) has produced and what actually will be executed:

```

SUB  #1,R2
ADD  R1,R2
BEZ  TAR
MUL  R3,R4
MOVE #10,R1
-----
TAR  -----
  
```

This instruction will be executed regardless of the condition.

This will be executed only if the branch has not been taken

94

This happens in the pipeline:Branch is taken

At this moment, both the condition (set by ADD) and the target address are known.



Penalty: 2 cycles

95

Branch is not take

At this moment the condition is known and the MOVE can go on.



Penalty: 1 cycle

96

What happens if the compiler is not able to find an instruction to be moved after the branch, into the branch delay slot?

97

In this case a NOP instruction (an instruction that does nothing) has to be placed after the branch. In this case the penalty will be the same as without delayed branching.

```

MUL  R2,R4
SUB  #1,R2
ADD  R1,R2
BEZ  TAR
NOP
MOVE #10,R1
-----
TAR  -----

```

Now, with R2, this instruction influences the following ones and cannot be moved from its place.

98

Branch Prediction

- In the last example we have considered that the *branch will not be taken* and we fetched the instruction following the branch; in the case the branch was taken the fetched instruction was discarded. As result, we had

branch penalty of

- 1 if the branch is **not** taken (prediction fulfilled)
- 2 if the branch is taken (prediction not fulfilled)

99

- Correct branch prediction is very important and can produce substantial performance improvements.
- Based on the predicted outcome, the respective instruction can be fetched, as well as the instructions following it, and they can be placed into the instruction queue.
- If, after the branch condition is computed, it turns out that the prediction was correct, execution continues.
- On the other hand, if the prediction is not fulfilled, the fetched instruction(s) must be discarded and the correct instruction must be fetched.

100

- To take full advantage of branch prediction, we can have the instructions not only fetched but also begin execution. **This is known as speculative execution.**
- Speculative execution means that instructions are executed before the processor is certain that they are in the correct execution path.** If it turns out that the prediction was correct, execution goes on without introducing any branch penalty.
- If, however, the prediction is not fulfilled, the instruction(s) started in advance and all their associated data must be purged and the state previous to their execution restored.

101

Branch prediction strategies:

- Static prediction
- Dynamic prediction

102

Static Branch Prediction

Static prediction techniques do not take into consideration execution history.

Static approaches:

1. Predict never taken (Motorola 68020): assumes that the branch is not taken.
2. Predict always taken: assumes that the branch is taken.
3. Predict depending on the branch direction (PowerPC 601):
 - predict branch taken for backward branches;
 - predict branch not taken for forward branches.

103

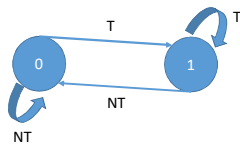
Dynamic Branch Prediction

- Dynamic prediction techniques improve the accuracy of the prediction by recording the history of conditional branches.
- **One-Bit Prediction Scheme**
- **Two-Bit Prediction Scheme**

104

One-Bit Prediction Scheme

- One-bit is used in order to record if the last execution resulted in a branch taken or not. The system predicts the same behavior as for the last time.



105

Dynamic Branch Prediction

When a branch is almost always taken, then when it is not taken, we will predict incorrectly twice, rather than once:

```

-----
LOOP -----
-----
BNZ  LOOP
-----
  
```

106

Dynamic Branch Prediction

- After the loop has been executed for the first time and left, it will be remembered that BNZ has not been taken. Now, when the loop is executed again, after the first iteration there will be a false prediction; following predictions are OK until the last iteration, when there will be a second false prediction.
- In this case the result is even worse than with static prediction considering that backward loops are always taken (PowerPC 601 approach).

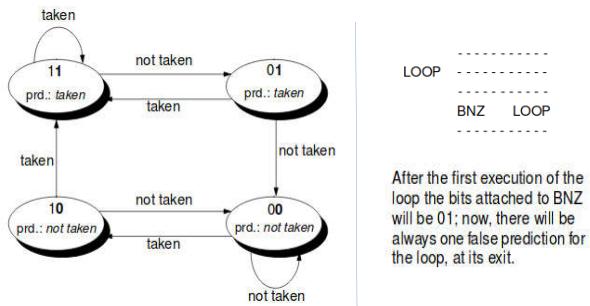
107

Two-Bit Prediction Scheme

- With a two-bit scheme predictions can be made depending on the last two instances of execution.
- A typical scheme is to change the prediction only if there have been two incorrect predictions in a row.

108

Two-Bit Prediction Scheme



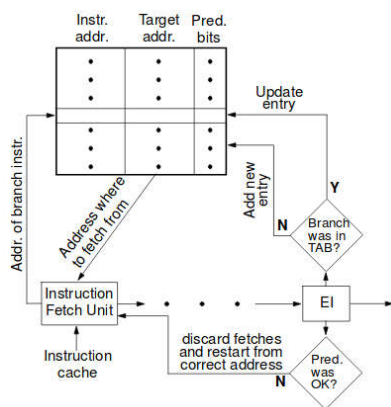
109

Branch History Table

- History information can be used not only to predict the outcome of a conditional branch but also to avoid recalculation of the target address.
- Together with the bits used for prediction, the target address can be stored for later use in a branch history table

110

Branch History Table



111

Branch History Table

- Address where to fetch from : If the branch instruction is not in the table the next instruction (address PC+1) is to be fetched. If the branch instruction is in the table first of all a prediction based on the prediction bits is made. Depending on the prediction outcome the next instruction (address PC+1) or the instruction at the target address is to be fetched.
- Update entry : If the branch instruction has been in the table, the respective entry has to be updated to reflect the correct or incorrect prediction.
- Add new entry: If the branch instruction has not been in the table, it is added to the table with the corresponding information concerning branch outcome and target address. If needed one of the existing table entries is discarded. Replacement algorithms similar to those for cache memories are used.

112

Branch History Table

- Using dynamic branch prediction with history tables up to 90% of predictions can be correct.
- Both Pentium and PowerPC 620 use speculative execution with dynamic branch prediction based on a branch history table

113

TERIMA KASIH

114